

ELEG 5040

Tutorial 1

Introduction to Python

Kai KANG (kkang@ee.cuhk.edu.hk)

Outline

- Tutorial Plan
- Piazza System
- Python Basics
- Popular Python Packages

Tutorial Plan

Week	Content	Note
1		
2	Python Introduction	
3	Theano	
4-5	CUDA/GPU Programming (Invited Talks)	Jan 24 & 25
6-7	Deep Learning Toolbox	
8-10	Caffe	
11-12	Research Experience on Deep Learning	
13-14	Review	

Piazza System

- Course website: www.piazza.com/cuhk.edu.hk/spring2015/eleg5040/home
- Resources: piazza.com/cuhk.edu.hk/spring2015/eleg5040/resources
 - Announcements (Email preferences)
 - Lecture and tutorial notes
 - Homework and solutions
 - Reading materials
- Q & A Section
 - Public and private questions, anonymous posts (if you are shy)
 - Collaborative editing: anyone can contribute
 - Image, attached files, LaTeX math forms, code highlighting, etc

Why Python?

Why Python?

If programming languages were vehicles



Python is great for everyday tasks: easy to drive, versatile, comes with all the conveniences built in. It isn't fast or sexy, but neither are your errands.

Why Python?

- Convenience
 - Modules
- Interpreted and interactive
- Clear and concise syntax
- Dynamic typing
- Portable
 - High-level
 - Object-oriented
 - Interfaces to many libraries
 - Large community

Online Materials

- Official Python tutorial (<https://docs.python.org/3/tutorial/index.html#tutorial-index>)
- CodeAcademy interactive tutorial (<http://www.codecademy.com/en/tracks/python>)
- Dive Into Python free digital book (<http://www.diveintopython.net>)
- Python Challenge (<http://www.pythonchallenge.com>)

Python Basics

Numbers

Arithmetics

```
>>> 2 + 2
4
>>> 50 - 5*6
20
>>> (50 - 5*6) / 4
5.0
>>> 8 / 5  # division always returns a floating point number
1.6
```

Floor division and remainder

```
>>> 17 / 3  # classic division returns a float
5.666666666666667
>>>
>>> 17 // 3  # floor division discards the fractional part
5
>>> 17 % 3  # the % operator returns the remainder of the division
2
>>> 5 * 3 + 2  # result * divisor + remainder
17
```

Numbers

Power operations

```
>>> 5 ** 2    # 5 squared
25
>>> 2 ** 7    # 2 to the power of 7
128
```

Variables and assignments

```
>>> width = 20
>>> height = 5 * 9
>>> width * height
900
```

Numbers

Beyond int and float

```
>>> from fractions import Fraction
>>> Fraction(16, -10)
Fraction(-8, 5)
>>> Fraction(123)
Fraction(123, 1)
>>> Fraction()
Fraction(0, 1)
>>> Fraction('3/7')
Fraction(3, 7)
```

```
>>> 3+5j
(3+5j)
>>> 3+5J
(3+5j)
>>> 6+7j + 8-2J
(14+5j)
>>> complex('6-5j')
(6-5j)
```

Strings

Use either single quotes or double quotes

```
>>> 'spam eggs'    # single quotes
'spam eggs'
>>> 'doesn\'t'     # use \' to escape the single quote...
'doesn't'
>>> "doesn't"      # ...or use double quotes instead
'doesn't'
>>> '"Yes," he said.'
'"Yes," he said.'
```

Use \ to escape sequences

```
>>> "\"Yes,\" he said."
'"Yes," he said.'
>>> '"Isn\'t," she said.'
'"Isn\'t," she said.'
>>> 'Hello, World!\n'
'Hello, World!\n'
>>> print('Hello, World!\n')
Hello, World!

>>>
```

Strings

Concatenation and repetition

```
>>> 'He' + 'llo'
'Hello'
>>> ('He' + 'llo') * 3
'HelloHelloHello'
>>> 'He' 'llo'
'Hello'
>>> ('He' "llo") * 3
'HelloHelloHello'
```

Indexing (zero-based)

```
>>> word = 'Python'
>>> word[0]    # character in position 0
'P'
>>> word[5]    # character in position 5
'n'
>>> word[-1]   # last character
'n'
>>> word[-2]   # second-last character
'o'
>>> word[0:2]  # characters from position 0 (included) to 2 (excluded)
'Py'
>>> word[2:5]  # characters from position 2 (included) to 5 (excluded)
'tho'
```

Lists

Lists may not contain items with same types

```
list1 = ['physics', 'chemistry', 1997, 2000];  
list2 = [1, 2, 3, 4, 5 ];  
list3 = ["a", "b", "c", "d"];
```

Accessing, updating and deleting items

```
>>> list1 = ['physics', 'chemistry', 1997, 2000];  
>>> list2 = [1, 2, 3, 4, 5, 6, 7 ];  
>>> print "list1[0]: ", list1[0]  
list1[0]: physics  
>>> print "list2[1:5]: ", list2[1:5]  
list2[1:5]: [2, 3, 4, 5]  
>>> list1[2]=2001  
>>> list1  
['physics', 'chemistry', 2001, 2000]  
>>> del list1[2]  
>>> list1  
['physics', 'chemistry', 2000]
```

Concatenating and repeating

```
>>> [1, 2, 3] + [4, 5]  
[1, 2, 3, 4, 5]  
>>> [6, 7] * 4  
[6, 7, 6, 7, 6, 7, 6, 7]
```

Tuples

Tuples are very similar to lists, only difference is that tuples are immutable

```
>>> tup1 = ('physics', 'chemistry', 1997, 2000);
>>> tup2 = (1, 2, 3, 4, 5, 6, 7 );
>>> print "tup1[0]: ", tup1[0]
tup1[0]: physics
>>> print "tup2[1:5]: ", tup2[1:5]
tup2[1:5]: (2, 3, 4, 5)
>>> tup1 = (50,)
```

No enclosing delimiters needed

```
>>> a, b = 0, 1
>>> a, b = b, a + b
>>> a, b
(1, 1)
```


Dictionaries

Creating dictionaries with {}

```
>>> dict = {'Alice': '2341', 'Beth': '9102', 'Cecil': '3258'}
>>> dict
{'Beth': '9102', 'Alice': '2341', 'Cecil': '3258'}
>>> dict1 = { 'abc': 456 };
>>> dict2 = { 'abc': 123, 98.6: 37 };
>>> dict2
{98.6: 37, 'abc': 123}
```

Indexing with keys

```
>>> dict['Alice']
'2341'
>>> dict1['abc']
456
>>> dict2[98.6]
37
```

Looping through keys

```
>>> for key in dict2:
...     print key, dict2[key]
...
98.6 37
abc 123
```

if Statements

Indentation is important in Python, semicolons are not

```
>>> a = 3
>>> b = 4
>>> if a < b:
...     print(a)
... else:
...     print(b)
...
3
>>> print(a if a < b else b)
3
```

No switch and case statements, use if..elif..elif.. sequence

```
>>> if x < 0:
...     x = 0
...     print('Negative changed to zero')
... elif x == 0:
...     print('Zero')
... elif x == 1:
...     print('Single')
... else:
...     print('More')
```

for Statements

for statements iterate within sequences

```
>>> # Measure some strings:
... words = ['cat', 'window', 'defenestrate']
>>> for w in words:
...     print(w, len(w))
...
cat 3
window 6
defenestrate 12
```

Looping in a range

```
>>> for i in range(5):
...     print i
0
1
2
3
4
>>> range(1, 5)
[1, 2, 3, 4]
>>> range(1, 5, 2)
[1, 3]
```

break and continue

break jumps out the smallest loop

```
>>> for n in range(2, 10):
...     for x in range(2, n):
...         if n % x == 0:
...             print n, 'equals', x, '*', n/x
...             break
...     else:
...         # loop fell through without finding a factor
...         print n, 'is a prime number'
```

continue jumps to next iteration

```
>>> for num in range(2, 10):
...     if num % 2 == 0:
...         print "Found an even number", num
...         continue
...     print "Found a number", num
```

pass Statements

pass statements do nothing but make syntax right

```
>>> while True:
...     pass    # Busy-wait for keyboard interrupt (Ctrl+C)
...
```

```
>>> class MyEmptyClass:
...     pass
...
```

```
>>> def initlog(*args):
...     pass    # Remember to implement this!
...
```

Functions

Define a function

```
>>> def fib(n):      # write Fibonacci series up to n
...     """Print a Fibonacci series up to n."""
...     a, b = 0, 1
...     while a < n:
...         print a,
...         a, b = b, a+b
...
>>> # Now call the function we just defined:
... fib(2000)
0 1 1 2 3 5 8 13 21 34 55 89 144 233 377 610 987 1597
```

Functions are objects

```
>>> fib
<function fib at 10042ed0>
>>> f = fib
>>> f(100)
0 1 1 2 3 5 8 13 21 34 55 89
```

Functions

return statements

```
>>> def fib2(n): # return Fibonacci series up to n
...     """Return a list containing the Fibonacci series up to
n."""
...     result = []
...     a, b = 0, 1
...     while a < n:
...         result.append(a)      # see below
...         a, b = b, a+b
...     return result
...
>>> f100 = fib2(100)      # call it
>>> f100                  # write the result
[0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89]
```

Functions

Default argument values

```
>>> def inc(a, b = 1):  
...     return a + b  
...  
>>> inc(10)  
11  
>>> inc(10, 5)  
15
```

Keyword arguments

```
>>> def inc(st, step=1):  
...     return st + step  
...  
>>> inc(10)  
11  
>>> inc(10, 5)  
15  
>>> inc(10, step=6)  
16
```

Arbitrary argument list

```
def write_multiple_items(file, separator, *args):  
    file.write(separator.join(args))
```


Classes

Define a class: instance variables and methods

```
class MyClass:
    """A simple example class"""
    i = 12345
    def f(self):
        return 'hello world'

>>> x = MyClass()
>>> x.i
12345
>>> x.f()
'hello world'
```

Instantiations: `__init__()`

```
>>> class Complex:
...     def __init__(self, realpart, imagpart):
...         self.r = realpart
...         self.i = imagpart
...
>>> x = Complex(3.0, -4.5)
>>> x.r, x.i
(3.0, -4.5)
```

Classes

Instance variables

```
>>> class Dog:
...     def __init__(self, name, kind='canine'):
...         self.name = name
...         self.kind = kind
...
>>> d = Dog('Fido')
>>> e = Dog('Buddy', 'husky')
>>> d.kind
'canine'
>>> e.kind
'husky'
>>> d.name
'Fido'
>>> e.name
'Buddy'
```

Classes

Inheritance

```
>>> class Dog:
...     def __init__(self, name, kind='canine'):
...         self.name = name
...         self.kind = kind
...
```

```
>>> class Husky(Dog):
...     def __init__(self, name):
...         Dog.__init__(self, name, 'husky')
...
```

```
>>> h = Husky('Rocky')
```

```
>>> h.name
```

```
'Rocky'
```

```
>>> h.kind
```

```
'husky'
```

Run Python Scripts

```
#!/usr/bin/python3

class Duck:
    def quack(self):
        print('Quaaack!')

    def walk(self):
        print('Walks like a duck.')

def main():
    donald = Duck()
    donald.quack()
    donald.walk()

if __name__ == "__main__": main()
```

Popular Packages

- cPickle, Protocol Buffers
- NumPy, SciPy
- scikit-learn, scikit-image, PIL, matplotlib
- h5py, leveldb
- ipython
- ...

NumPy for Matlab Users

- http://wiki.scipy.org/NumPy_for_Matlab_Users

MATLAB	numpy	Notes
<code>a && b</code>	<code>a and b</code>	short-circuiting logical AND operator (Python native operator); scalar arguments only
<code>a b</code>	<code>a or b</code>	short-circuiting logical OR operator (Python native operator); scalar arguments only
<code>1*i, 1*j, 1i, 1j</code>	<code>1j</code>	complex numbers
<code>eps</code>	<code>spacing(1)</code>	Distance between 1 and the nearest floating point number
<code>ode45</code>	<code>scipy.integrate.ode(f).set_integrator('dopri5')</code>	integrate an ODE with Runge-Kutta 4,5
<code>ode15s</code>	<code>scipy.integrate.ode(f).\nset_integrator('vode', method='bdf', order=15)</code>	integrate an ODE with BDF

NumPy for Matlab Users

- Linear Algebra Equivalents

MATLAB	numpy.array	numpy.matrix
<code>ndims(a)</code>	<code>ndim(a)</code> or <code>a.ndim</code>	
<code>numel(a)</code>	<code>size(a)</code> or <code>a.size</code>	
<code>size(a)</code>	<code>shape(a)</code> or <code>a.shape</code>	
<code>size(a,n)</code>	<code>a.shape[n-1]</code>	
<code>[1 2 3; 4 5 6]</code>	<code>array([[1.,2.,3.], [4.,5.,6.]])</code>	<code>mat([[1.,2.,3.], [4.,5.,6.]])</code> or
<code>[a b; c d]</code>	<code>vstack([hstack([a,b]), hstack([c,d])])</code>	<code>bmat('a b; c d')</code>
<code>a(end)</code>	<code>a[-1]</code>	<code>a[:,-1][0,0]</code>
<code>a(2,5)</code>	<code>a[1,4]</code>	
<code>a(2,:)</code>	<code>a[1]</code> or <code>a[1,:]</code>	
<code>a(1:5,:)</code>	<code>a[0:5]</code> or <code>a[:5]</code> or <code>a[0:5,:]</code>	
<code>a(end-4:end,:)</code>	<code>a[-5:]</code>	
<code>a(1:3,5:9)</code>	<code>a[0:3][:,4:9]</code>	
<code>a([2,4,5],[1,3])</code>	<code>a[ix_([1,3,4],[0,2])]</code>	

NumPy for Matlab Users

- Linear Algebra Equivalents

MATLAB	numpy.array	numpy.matrix
<code>a(3:2:21,:)</code>	<code>a[2:21:2,:]</code>	
<code>a(1:2:end,:)</code>	<code>a[::2,:]</code>	
<code>a(end:-1:1,:) or flipud(a)</code>	<code>a[::-1,:]</code>	
<code>a([1:end 1],:)</code>	<code>a[r_[:len(a),0]]</code>	
<code>a.'</code>	<code>a.transpose()</code> or <code>a.T</code>	
<code>a'</code>	<code>a.conj().transpose()</code> or <code>a.conj().T</code>	<code>a.H</code>
<code>a * b</code>	<code>dot(a,b)</code>	<code>a * b</code>
<code>a .* b</code>	<code>a * b</code>	<code>multiply(a,b)</code>
<code>a./b</code>	<code>a/b</code>	
<code>a.^3</code>	<code>a**3</code>	<code>power(a,3)</code>
<code>(a>0.5)</code>	<code>(a>0.5)</code>	
<code>find(a>0.5)</code>	<code>nonzero(a>0.5)</code>	
<code>a(:,find(v>0.5))</code>	<code>a[:,nonzero(v>0.5)[0]]</code>	<code>a[:,nonzero(v.A>0.5)[0]]</code>
<code>a(:,find(v>0.5))</code>	<code>a[:,v.T>0.5]</code>	<code>a[:,v.T>0.5]</code>

NumPy for Matlab Users

- Linear Algebra Equivalents

MATLAB	numpy.array	numpy.matrix
<code>a(a<0.5)=0</code>	<code>a[a<0.5]=0</code>	
<code>a .* (a>0.5)</code>	<code>a * (a>0.5)</code>	<code>mat(a.A * (a>0.5).A)</code>
<code>a(:) = 3</code>	<code>a[:] = 3</code>	
<code>y=x</code>	<code>y = x.copy()</code>	
<code>y=x(2,:)</code>	<code>y = x[1,:].copy()</code>	
<code>y=x(:)</code>	<code>y = x.flatten(1)</code>	
<code>1:10</code>	<code>arange(1.,11.)</code> or <code>r_[1.:11.]</code> or <code>r_[1:10:10j]</code>	<code>mat(arange(1.,11.))</code> or <code>r_[1.:11.,'r']</code>
<code>0:9</code>	<code>arange(10.)</code> or <code>r_[0:10.]</code> or <code>r_[0:9:10j]</code>	<code>mat(arange(10.))</code> or <code>r_[0:10.,'r']</code>
<code>[1:10]'</code>	<code>arange(1.,11.)[:, newaxis]</code>	<code>r_[1.:11.,'c']</code>
<code>zeros(3,4)</code>	<code>zeros((3,4))</code>	<code>mat(...)</code>
<code>zeros(3,4,5)</code>	<code>zeros((3,4,5))</code>	<code>mat(...)</code>
<code>ones(3,4)</code>	<code>ones((3,4))</code>	<code>mat(...)</code>
<code>eye(3)</code>	<code>eye(3)</code>	<code>mat(...)</code>
<code>diag(a)</code>	<code>diag(a)</code>	<code>mat(...)</code>
<code>diag(a,0)</code>	<code>diag(a,0)</code>	<code>mat(...)</code>
<code>rand(3,4)</code>	<code>random.rand(3,4)</code>	<code>mat(...)</code>

Reading and Writing Mat Files

```
>>> import scipy.io as sio
>>> import numpy as np
>>> x = np.arange(12).reshape((3,4))
>>> x
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>> sio.savemat('example.mat', {'x': x})
>>> y = sio.loadmat('example.mat')
>>> y
{'x': array([[ 0,  1,  2,  3],
             [ 4,  5,  6,  7],
             [ 8,  9, 10, 11]]), '__version__': '1.0', '__header__':
'MATLAB 5.0 MAT-file Platform: posix, Created on: Fri Jan 16 13:04:04
2015', '__globals__': []}
>>> new_x = y['x']
>>> new_x
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
```

Pickle

```
>>> import cPickle
>>> import numpy as np
>>> x = np.arange(12).reshape((3,4))
>>> y = np.random.rand(3,4)
```

Serialization

```
>>> pkl_file = open('example.pkl', 'wb')
>>> cPickle.dump((x,y), pkl_file)
>>> pkl_file.close()
```

De-serialization

```
>>> new_x, new_y = cPickle.load(open('example.pkl', 'rb'))
>>> new_x
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
>>> new_y
array([[ 0.64956048,  0.79727908,  0.28350831,  0.07649368],
       [ 0.11551531,  0.11434357,  0.89143651,  0.23523877],
       [ 0.21506583,  0.16611042,  0.15059062,  0.30306736]])
```

References

- Official Python tutorial (<https://docs.python.org/3/tutorial/index.html#tutorial-index>)
- CodeAcademy interactive tutorial (<http://www.codecademy.com/en/tracks/python>)
- Dive Into Python free digital book (<http://www.diveintopython.net>)
- Python Challenge (<http://www.pythonchallenge.com>)
- NumPy for Matlab Users (http://wiki.scipy.org/NumPy_for_Matlab_Users)
- Scikit-learn (<http://scikit-learn.org/stable/>)
- SciPy, NumPy, matplotlib (<http://www.scipy.org>)
- h5py (<http://docs.h5py.org/en/latest/>)
- Style Guide for Python Code (<https://www.python.org/dev/peps/pep-0008/>)
- Google Python Style Guide (<https://google-styleguide.googlecode.com/svn/trunk/pyguide.html>)

Tutorial 2

Introduction to Theano

Kai KANG (kkang@ee.cuhk.edu.hk)

Outline

- CUDA/GPU Programming Invited Talks
- Theano Basics
- Beyond Basics
- Theano in Machine Learning and Deep Learning

CUDA/GPU Programming Tutorials

- Speaker: Bin Zhou, PhD
 - NVIDIA CUDA Fellow, USTC Adjunct Research Professor
 - Chief Scientist and Director of Marine Remote Sensing & Information Processing Lab, SDIOI
- First tutorial: 9:30 - 12:15 on Saturday, Jan 24, 2015, LSB LT6
- Second tutorial: 9:30 - 12:15 on Sunday, Jan 25, 2015, TY Wong Hall at SHB

Theano Basics

What is Theano?

- Symbolic computation library
- CPU and GPU infrastructure
- Optimized compiler

Online Materials

- Theano (<http://deeplearning.net/software/theano/index.html>)
- Theano introduction (<http://deeplearning.net/software/theano/introduction.html#introduction>)
- Installation guides (<http://deeplearning.net/software/theano/install.html#install>)
- Theano tutorials (<http://deeplearning.net/software/theano/tutorial/index.html#tutorial>)
- Documentation (<http://deeplearning.net/software/theano/library/index.html#libdoc>)
- GitHub page (<https://github.com/Theano/Theano>)

Algebra Basics

```
import theano
```

```
>>> import theano
Using gpu device 0: GeForce GTX 670
>>> import theano.tensor as T
>>> from theano import function
```

```
define variables
```

```
>>> x = T.fscalar('x')
>>> y = T.fscalar('y')
>>> x.type
TensorType(float32, scalar)
>>> y.type
TensorType(float32, scalar)
```

Algebra Basics

symbolic operations

```
>>> z = x + y
>>> z.type
TensorType(float32, scalar)
>>> w = x * y
>>> s = z + w
```

pretty print expressions

```
>>> from theano import pp
>>> pp(z)
'(x + y)'
>>> pp(w)
'(x * y)'
>>> s = z + w
>>> pp(s)
'((x + y) + (x * y))'
```

Algebra Basics

compile into functions

```
>>> from theano import function
>>> f1 = function([x, y], z)
>>> f2 = function([x, y], w)
>>> f3 = function([x, y], s)
```

call functions to get NumPy array

```
>>> f1(1.0, 2.0)
array(3.0, dtype=float32)
>>> f2(1.0, 2.0)
array(2.0, dtype=float32)
>>> f3(1.0, 2.0)
array(5.0, dtype=float32)
```

Algebra Basics

- 3 Steps in using Theano
 - Step 1: define symbolic variables

```
>>> x = T.dscalar('x')
>>> y = T.dscalar('y')
```
 - Step 2: define symbolic expressions using defined variables

```
>>> z = x + y
```
 - Step 3: compile expressions into functions

```
>>> f = function([x, y], z)
```

Example

Add two matrices

```
>>> m1 = T.dmatrix()
>>> m2 = T.dmatrix()
>>> sum = m1 + m2
>>> msum = function([m1, m2], sum)

>>> import numpy as np
>>> msum([[1, 3, 5], [7, 9, 11]], [[2, 4, 6], [8, 10, 12]])
array([[ 3.,  7., 11.],
       [15., 19., 23.]])
>>> msum(np.random.rand(2, 3), np.random.rand(2, 3))
array([[ 1.07868464,  1.01885801,  0.69778457],
       [ 0.9562832 ,  1.56608957,  1.19840735]])
>>> msum(np.random.rand(2, 3), np.random.rand(2, 3))
array([[ 1.04984599,  0.42415405,  1.51894434],
       [ 0.80635964,  0.17668743,  1.72312159]])
```

Data Types

- **byte:** bscalar, bvector, bmatrix, brow, bcol, btensor3, btensor4
- **16-bit integers:** wscalar, wvector, wmatrix, wrow, wcol, wtensor3, wtensor4
- **32-bit integers:** iscalar, ivector, imatrix, irow, icol, itensor3, itensor4
- **64-bit integers:** lscalar, lvector, lmatrix, lrow, lcol, ltensor3, ltensor4
- **float:** fscalar, fvector, fmatrix, frow, fcol, ftensor3, ftensor4
- **double:** dscalar, dvector, dmatrix, drow, dcol, dtensor3, dtensor4
- **complex:** cscalar, cvector, cmatrix, crow, ccol, ctensor3, ctensor4


Beyond Basics

Elementwise Operations

```
>>> x = T.dmatrix('x')
>>> s = 1 / (1 + T.exp(-x))
>>> t = T.tanh(x)
>>> sigmoid = function([x], s)
>>> tanh = function([x], t)
>>> sigmoid([[0, 1], [-1, -2]])
array([[ 0.5          ,  0.73105858],
       [ 0.26894142,  0.11920292]])
>>> tanh([[0, 1], [-1, -2]])
array([[ 0.          ,  0.76159416],
       [-0.76159416, -0.96402758]])
>>> 2 * sigmoid([[0, 2], [-2, -4]]) - 1
array([[ 0.          ,  0.76159416],
       [-0.76159416, -0.96402758]])
```

Many mathematical operations, such as `exp`, `log`, `abs`, `sin`, `cos`, `+`, `-`, `*`, `/`, are elementwise operations

Elementwise Operations



```
Plot[{LogisticSigmoid[x],  
      Tanh[x],  
      2*LogisticSigmoid[2 x]-1},  
      {x, -2, 2}]
```

[WolframAlpha](https://www.wolframalpha.com)

Multiple Outputs

Theano functions support multiple outputs in a list

```
>>> a, b = T.dmatrices('a', 'b')
>>> diff = a - b
>>> abs_diff = abs(diff)
>>> squared_diff = diff ** 2
>>> f = function([a, b], [diff, abs_diff, squared_diff])
>>> f([[1, 1], [1, 1]], [[0, 1], [2, 3]])
[array([[ 1.,  0.],
        [-1., -2.])),
 array([[ 1.,  0.],
        [ 1.,  2.])),
 array([[ 1.,  0.],
        [ 1.,  4.]])]
>>> c, d, e = f([[1, 1], [1, 1]], [[0, 1], [2, 3]])
>>> c
array([[ 1.,  0.],
        [-1., -2.]])
>>> d
array([[ 1.,  0.],
        [ 1.,  2.]])
>>> e
array([[ 1.,  0.],
        [ 1.,  4.]])
```

Argument Default Values

Use Param class to specify function parameters

```
>>> from theano import Param
>>> input, step = T.dscalars('input', 'step')
>>> output = input + step
>>> incre = function([input, Param(step, default=1)], output)
>>> incre(33)
array(34.0)
>>> incre(33, 5)
array(38.0)
```

Set parameter values by name

```
>>> m = T.dscalar('m')
>>> output2 = output * m
>>> incre2 = function([input, Param(step, default=1), Param(m,
default=1, name='multiplier')], output2)
>>> incre2(33)
array(34.0)
>>> incre2(33, 5)
array(38.0)
>>> incre2(33, multiplier=2)
array(68.0)
```

Shared Variables

Theano functions can have internal states

```
>>> state = theano.shared(0.)
>>> inc_new = function([Param(step, default=1)], state+step,
updates=[(state, state+step)])
>>> inc_new()
array(1.0)
>>> inc_new()
array(2.0)
>>> inc_new(5)
array(7.0)
>>> inc_new(3)
array(10.0)
>>> inc_new()
array(11.0)
```

'updates' parameter is a list of shared variable and new expression pairs, which set the values of shared variables with those of new expressions

Shared Variables

getters and setters of shared variables

```
>>> state.get_value()  
array(11.0)  
>>> inc_new()  
array(12.0)  
>>> state.get_value()  
array(12.0)  
>>> inc_new(5)  
array(17.0)  
>>> state.get_value()  
array(17.0)  
>>> state.set_value(0)  
>>> inc_new()  
array(1.0)  
>>> state.get_value()  
array(1.0)  
>>> inc_new(5)  
array(6.0)  
>>> state.get_value()  
array(6.0)
```

Shared Variables

Shared variables are shared across functions

```
>>> dec = function([Param(step, default=1)], state-step,  
updates=[(state, state-step)])  
>>> dec()  
array(5.0)  
>>> dec(5)  
array(0.0)  
>>> state.get_value()  
array(0.0)  
>>> inc_new()  
array(1.0)  
>>> state.get_value()  
array(1.0)
```


Expression Replacement

Use ``givens`` parameter to replace part of an expression with a new expression

```
>>> x = T.dmatrix('x')
>>> s = 1 / (1 + T.exp(-x))
>>> t = T.tanh(x)
>>> tanh1 = function([x], t)
>>> tanh2 = function([y], 2 * s - 1, givens={x: 2 * y})
>>> tanh1([[0, 1], [3, 5]])
array([[ 0.          ,  0.76159416],
       [ 0.99505475,  0.9999092 ]])
>>> tanh2([[0, 1], [3, 5]])
array([[ 0.          ,  0.76159416],
       [ 0.99505475,  0.9999092 ]])
```

Use `theano.clone` to do more advanced replacements

```
>>> tanh3 = function([x], theano.clone(2 * s - 1, replace={x: 2 *
x}))
>>> tanh3([[0, 1], [3, 5]])
array([[ 0.          ,  0.76159416],
       [ 0.99505475,  0.9999092 ]])
```

Derivatives

Automatic derivatives with ``tensor.grad``

```
>>> x = T.dscalar('x')
>>> y = x ** 2
>>> dy = T.grad(y, wrt=x)
>>> f = function([x], dy)
>>> f(3.0)
array(6.0)
>>> f(5.0)
array(10.0)
```

Theano optimizes expressions when compiling functions

```
>>> pp(dy)
'((fill((x ** TensorConstant{2}), TensorConstant{1.0}) *
TensorConstant{2}) * (x ** (TensorConstant{2} -
TensorConstant{1}))))'
>>> pp(f.maker.fgraph.outputs[0])
'(TensorConstant{2.0} * x)'
```

Example - Logistic Regression

```
import numpy
import theano
import theano.tensor as T
rng = numpy.random

N = 400
feats = 784
D = (rng.randn(N, feats), rng.randint(size=N, low=0, high=2))
training_steps = 10000

# Declare Theano symbolic variables
x = T.matrix("x")
y = T.vector("y")
w = theano.shared(rng.randn(feats), name="w")
b = theano.shared(0., name="b")
print "Initial model:"
print w.get_value(), b.get_value()

# Construct Theano expression graph
p_1 = 1 / (1 + T.exp(-T.dot(x, w) - b)) # Probability that target = 1
prediction = p_1 > 0.5 # The prediction thresholded
xent = -y * T.log(p_1) - (1-y) * T.log(1-p_1) # Cross-entropy loss function
cost = xent.mean() + 0.01 * (w ** 2).sum() # The cost to minimize
gw, gb = T.grad(cost, [w, b]) # Compute the gradient of the cost
```

Example - Logistic Regression

```
# Compile
train = theano.function(
    inputs=[x,y],
    outputs=[prediction, xent],
    updates=((w, w - 0.1 * gw), (b, b - 0.1 * gb)))
predict = theano.function(inputs=[x], outputs=prediction)

# Train
for i in range(training_steps):
    pred, err = train(D[0], D[1])

print "Final model:"
print w.get_value(), b.get_value()
print "target values for D:", D[1]
print "prediction on D:", predict(D[0])
```

Theano in ML and DL

- Deep learning tutorial based on Theano (<http://deeplearning.net/tutorial/contents.html>)
 - Getting started tutorial (<http://deeplearning.net/tutorial/gettingstarted.html>)
- Pylearn2: machine learning library based on Theano (<http://deeplearning.net/software/pylearn2/>)
- GroundHog: RNN framework based on Theano (<https://github.com/lisa-groundhog/GroundHog>)
- Faster convolution methods (<https://github.com/Theano/Theano/wiki/Convolution-methods>)
- ...

Tutorial 3. Deep Learning Toolbox

Xingyu ZENG(xyzeng@ee.cuhk.edu.hk)

DeepLearnToolbox

- ▶ A open-source Matlab toolbox for Deep Learning

- ▶ You can download in

<https://github.com/rasmusbergpalm/DeepLearnToolbox>

- ▶ If you use this toolbox in your research please cite

@MASTERSTHESIS\{IMM2012-06284, author = "R. B. Palm", title = "Prediction as a candidate for learning deep hierarchical models of data", year = "2012", }



DeepLearnToolbox

▶ Advantage

- ▶ Matlab, easy to use
- ▶ Open-source

▶ Disadvantage

- ▶ Only CPU version, slow

Install Steps

1. Download the toolbox,
2. `Addpath(genpath('DeepLearnToolbox'))`



DeepLearnToolbox

- ▶ A Matlab toolbox for Deep Learning
 - ▶ NN/ - A library for Feedforward Backpropagation Neural Networks
 - ▶ CNN/ - A library for Convolutional Neural Networks
 - ▶ DBN/ - A library for Deep Belief Networks
 - ▶ SAE/ - A library for Stacked Auto-Encoders
 - ▶ CAE/ - A library for Convolutional Auto-Encoders
 - ▶ util/ - Utility functions used by the libraries
 - ▶ data/ - Data used by the examples
 - ▶ tests/ - unit tests to verify toolbox is working



DeepLearnToolbox

Feedforward Backpropagation Neural Networks

➤ Common Function

- `nnsetup.m`
 - To setup one network
- `nntrain.m`
 - To train one network
- `nnpredict.m`
 - To test samples with one network



DeepLearnToolbox

► nnsetup.m

Usage example:


```
nn = nnsetup([784 100 10]); % to build up one three layers network
```

```
nn.activation_function = 'sigm';
```

```
nn.output = 'softmax';
```

```
h=sigmoid(W1*x);
```

```
y=softmax(W2*h);
```



DeepLearnToolbox

► nntrain.m

Usage example:

```
nn.learningRate = 0.1;  
opts.numepochs = 1; %Number of full sweeps through data  
opts.batchsize = 100; %Take a mean gradient step over this many  
                    %samples  
  
nn = nntrain(nn, train_x, train_y, opts);
```



DeepLearnToolbox

► `nnpredict.m`


Usage example:

```
labels = nnpredict(nn, test_x);
```

Notes:

1. `labels`, the classes predicted by `nn`
2. `nn.a{end}`, the values of the output layer

```
[~,labels]=max(nn.a{end},[],2);
```



DeepLearnToolbox

Convolutional Neural Networks

➤ Common Function

➤ `cnnsetup.m`

- To setup one convolutional network

➤ `cnntrain.m`

- To train one convolutional network

➤ `cnnff.m`

- Forward step with one convolutional network


➤ `cnntest.m`

- To test samples with one convolutional network



DeepLearnToolbox

- Suppose input size: $152 \times 76 \times 3$



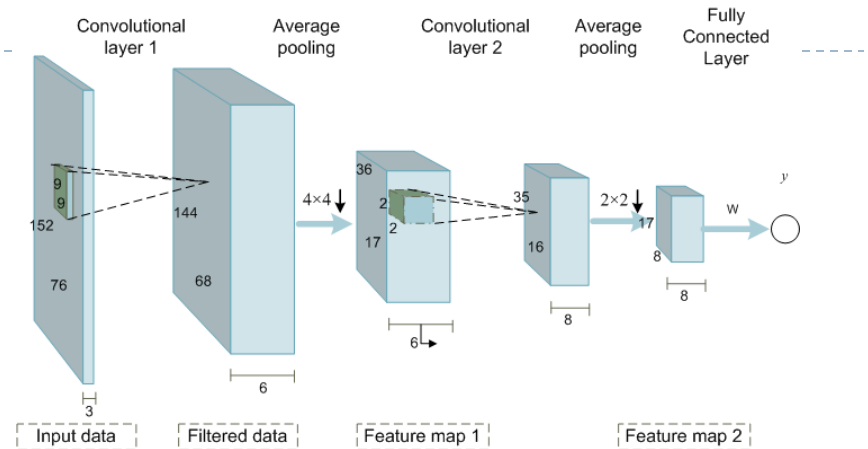
DeepLearnToolbox

► `cnnsetup.m`

Usage example:

```
cnn.layers = {  
    struct('type', 'i') %input layer  
    struct('type', 'c', 'outputmaps', 6, 'kernelsize', 9) %convolution layer  
    struct('type', 's', 'scale', 4) %sub sampling layer  
    struct('type', 'c', 'outputmaps', 8, 'kernelsize', 2) %convolution layer  
    struct('type', 's', 'scale', 2) %subsampling layer  
};
```

% the size of {train_x, train_y} is useful for setup cnn
cnn = cnnsetup(cnn, train_x, train_y);



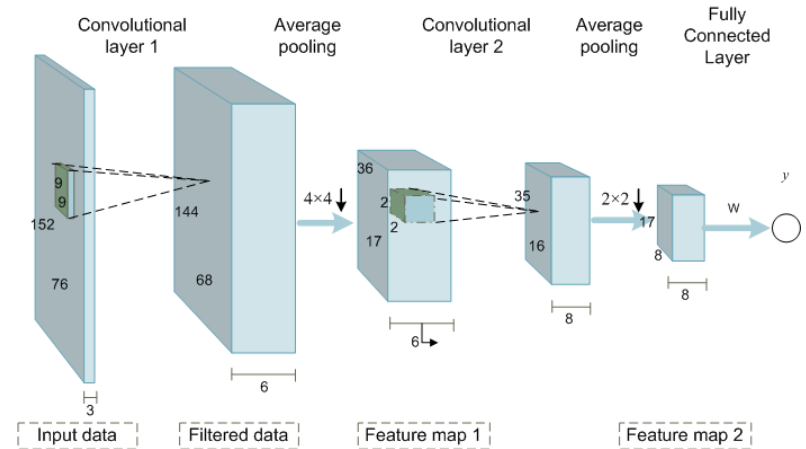
DeepLearnToolbox

► `cnntrain.m`

Usage example:

```
opts.alpha = 1; % learning rate
opts.batchsize = 50;
opts.numepochs = 1;
```

```
cnn = cnntrain(cnn, train_x, train_y, opts);
```

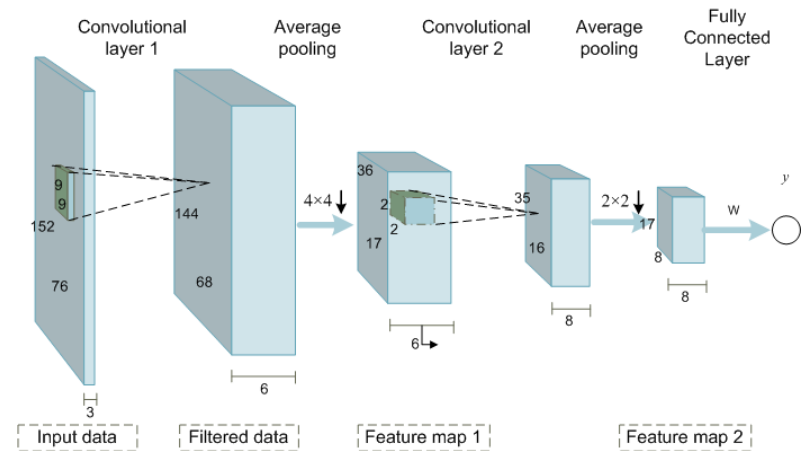


DeepLearnToolbox

► cnnff.m

Usage example:

```
cnn = cnnff(cnn, x);
```



Notes:

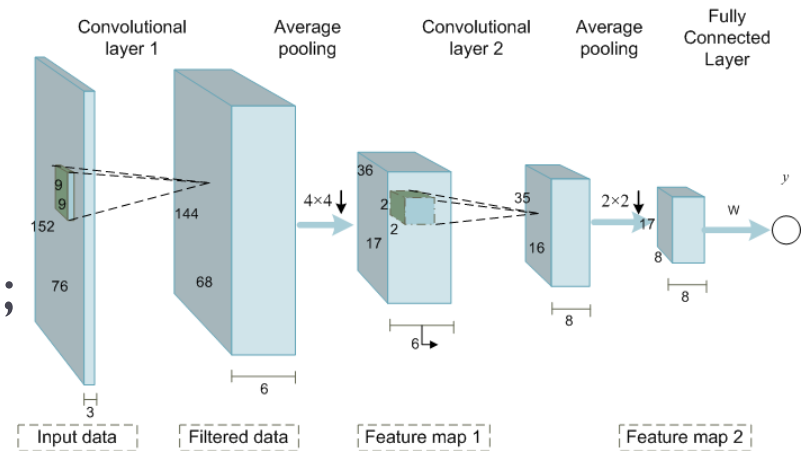
1. `cnn.o`, the output values of the output layer
2. `cnn.fv`, the feature of the input samples
3. `cnn.o=sigm(cnn.ffW * cnn.fv + repmat(cnn.ffb, 1, size(cnn.fv, 2)))`;

DeepLearnToolbox

► cnntest.m

Usage example:

```
[er, bad] = cnntest(cnn, test_x, test_y);
```



Notes:

1. `er`, error fraction value
2. `bad`, index of misclassified testing samples

DeepLearnToolbox

- ▶ More examples can be found in

<https://github.com/rasmusbergpalm/DeepLearnToolbox/blob/master/README.md>

- ▶ More details about deep learning,
Book, 'Learning Deep Architectures for AI'



Tutorial 4. MatConvNet

Xingyu ZENG(xyzeng@ee.cuhk.edu.hk)

MatConvNet

- ▶ A open-source Matlab toolbox for Convolution Network
- ▶ You can download in
<https://github.com/vlfeat/matconvnet>
- ▶ *If you use MatConvNet in your work, please cite:*
"MatConvNet - Convolutional Neural Networks for MATLAB", A. Vedaldi and K. Lenc, arXiv:1412.4564, 2014.



MatConvNet

- ▶ **Advantage**

- ▶ Matlab, easy to use
- ▶ Pretrained models(VGG,AlexNet)
- ▶ Support GPU

- ▶ **Disadvantage**

- ▶ Complicated than DeepLearnToolbox
- ▶ Support only Convolution Network



MatConvNet

- ▶ Installing and compiling the library
 - ▶ Matlab
 - ▶ run `<path to MatConvNet>/matlab/vl_setupnn`
 - ▶ `vl_compilenn();`
 - ▶ `vl_compilenn('enableGpu', true); % for gpu`
 - ▶ Shell
 - ▶ make `ARCH=<your arch> MATLABROOT=<path to MATLAB>`
 - ▶ make `ENABLE_GPU=y ARCH=<your arch> MATLABROOT=<path to MATLAB> CUDAROOT=<path to CUDA>`

Notes: the gpu version require at least MATLAB2013



MatConvNet

► Overview

1. Core functions

- Different type of layers, including convolution, dropout

2. Simple CNN functions

- A simple wrapper for CNN, including training, testing, display

3. Utility functions

- Some helper function to initialize the toolbox



MatConvNet

► Pretrained Models (VGG,AlexNet)

Usage:

1. Download a pre-trained CNN
2. `net=load('cnn.mat');`
3. `im=im-net.normalization.averageImage;`
4. `res=vl_simplenn(net,im);`

Notes:

1. `res(i).x`, output for i-th layer
2. Each model has one `averageImage` value



MatConvNet

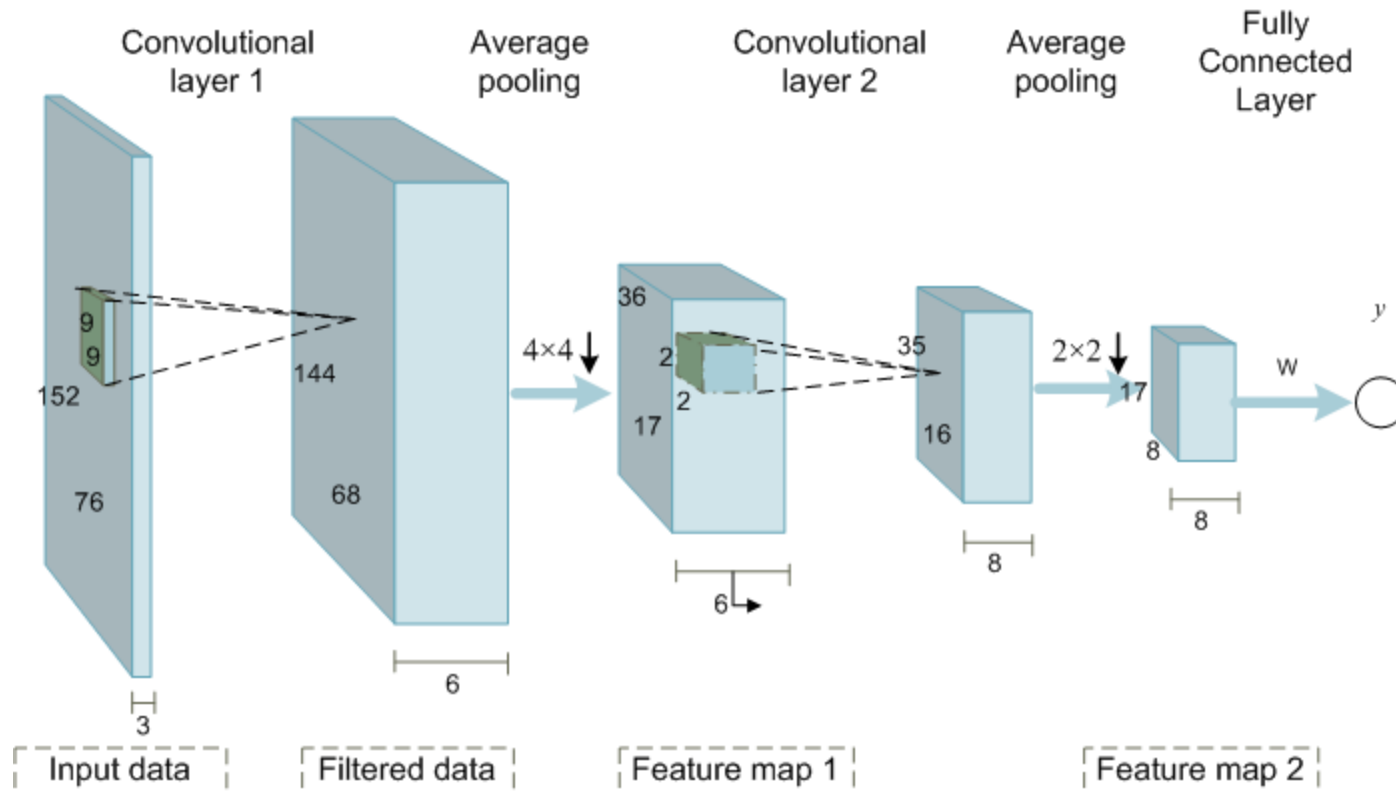
► Notes:

1. Format: Height*Width*Channels*Num
2. Pretrained Models:
 - VGG models
`Very Deep Convolutional Networks for Large-Scale Image Recognition', Karen Simonyan and Andrew Zisserman, arXiv technical report, 2014
 - AlexNet
`ImageNet Classification with Deep Convolutional Neural Networks', Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton, NIPS, 2012



MatConvNet

- Suppose input size: $152 \times 76 \times 3$



MatConvNet

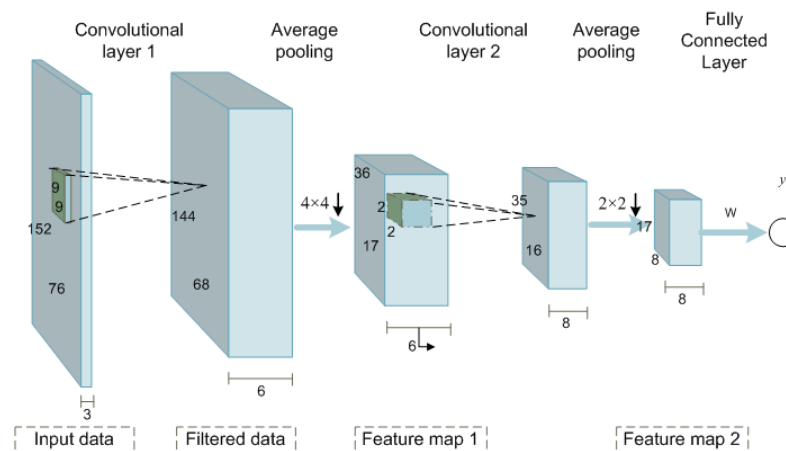
► Construct one CNN

```
net.layers = {} ;
```

```
net.layers{end+1} = struct('type','conv', ...  
    'filters', f*randn(9,9,3,6, 'single'), ...  
    'biases', zeros(1, 6, 'single'), ...  
    'stride', 1, ...  
    'pad', 0) ;  
  
net.layers{end+1} = struct('type','pool', ...  
    'method','avg', ...  
    'pool', [4 4], ...  
    'stride', 4, ...  
    'pad', 0) ;
```

.....

```
net.layers{end+1} = struct('type','softmaxloss') ;
```



MatConvNet

► Layers

1. Convolution Layer
2. Pooling Layer
3. RELU Layer
4. Dropout Layer
5. Softmax Layer
6. Log-loss layer
7. Softmax-log-loss-layer
8. Custom Layer
 1. `layer.type = 'custom'`
 2. `layer.forward`: a function handle computing the block.
 3. `layer.backward`: a function handle computing the block derivative.



MatConvNet

► Training one CNN

Usage:

1. `opts.train.batchSize=100;`
2. `opts.train.numEpochs=100;`
3. `opts.train.useGpu=false;`
4. `opts.train.learningRate=0.001;`
5. `[net,info]=cnn_train(net,imdb,@getBatch,opts);`

Notes:

1. `getBatch`: one function, `[im,labels]=getBatch(imdb,ind);`
2. `info`: error, objective, topFiveError....



MatConvNet

- ▶ More details can be found in

<http://www.vlfeat.org/matconvnet/matconvnet-manual.pdf>

- ▶ More examples can be found in

<https://github.com/vlfeat/matconvnet/tree/master/examples>

- ▶ The homepage of MatConvNet:

<http://www.vlfeat.org/matconvnet/>



Tutorial 7

Introduction to Caffe

Kai KANG (kkang@ee.cuhk.edu.hk)

Outline

- General Introduction
 - What is Caffe?
 - Why using Caffe?
 - How to get started?
- Caffe in details
 - Model and protocol buffers
 - Forward and backward computations
 - Loss and supervision
 - Solver and optimizations
 - Data and database
 - Interfaces and tools

What is Caffe?

- Open source deep learning framework maintained by Berkeley Vision and Learning Center (BVLC)
- Created by Yangqing Jia as an improved version of DeCaf
- Mainly written in C++ and CUDA C with Python and Matlab interfaces
- Tools, reference models and sample recipes included

Why Using Caffe? (Official)

- Expression: models and optimizations are defined as plaintext schemas instead of code.
- Speed: for research and industry alike speed is crucial for state-of-the-art models and massive data.
- Modularity: new tasks and settings require flexibility and extension.
- Openness: scientific and applied progress call for common code, reference models, and reproducibility.
- Community: academic research, startup prototypes, and industrial applications all share strength by joint discussion and development in a BSD-2 project.

Why We Use Caffe?

- Good starting point
- Reliability, especially for large scale problem
- Speed
- Popularity

How to Get Started?

- Official website (<http://caffe.berkeleyvision.org>)
- Download from the GitHub page (<https://github.com/BVLC/caffe>)
- Follow the installation instructions to compile and test (<http://caffe.berkeleyvision.org/installation.html>)
- Try the tutorials and reference models (<http://caffe.berkeleyvision.org/tutorial/>)
- Look through the detailed API documentations (<http://caffe.berkeleyvision.org/doxygen/annotated.html>)

Model and Protocol Buffers

What is a Network?


- Text representations: human-readable network configurations
- Parameter representations: trained network parameters
- Class representations: network data and behavior implementations

What is a Caffe net?

```
name: "CaffeNet"
layers {
  name: "data"
  type: IMAGE_DATA
  ...
}
layers {
  name: "conv1"
  type: CONVOLUTION
  ...
}
...

layers {
  name: "loss"
  type: SOFTMAX_LOSS
  ...
}
```

```
message NetParameter {
  optional string name = 1;
  repeated LayerParameter layers =
    2;
  repeated string input = 3;
  repeated int32 input_dim = 4;
  optional bool force_backward = 5
    [default = false];
  optional NetState state = 6;
}
```



Text
Representations



Net Definitions

Parameter
Representations



Data Structure

Class
Representations



Implementation

Protocol Buffers

- Google open-source project
- Represent structured data
- Support serialization and de-serialization
- Include C++, Python and Java interfaces
- Compilers compile protocol buffers into classes to access the data

Example - Blobs

- A Blob is a wrapper over the actual data in Caffe
- It's a 4D tensor with dimensions of (num, channels, height and width)
- It can represent data, results and network parameters
 - Data: $20 * 3 * 256 * 256$
 - Results: $20 * 96 * 128 * 128$
 - Convolution kernels: $96 * 3 * 7 * 7$
 - Fully-connected vector: $4096 * 128 * 1 * 1$
- Blobs store results of both forward and backward propagations

Example - Blobs

Protocol Buffers definitions

```
message BlobProto {  
    optional int32 num = 1 [default = 0];  
    optional int32 channels = 2 [default = 0];  
    optional int32 height = 3 [default = 0];  
    optional int32 width = 4 [default = 0];  
    repeated float data = 5 [packed = true];  
    repeated float diff = 6 [packed = true];  
}
```

Represent multiple blobs

```
message BlobProtoVector {  
    repeated BlobProto blobs = 1;  
}
```

Example - Blobs

C++ class definitions

```
class Blob {
public:
    Blob()
        : data_(), diff_(), num_(0), channels_(0), height_(0), width_(0),
          count_(0), capacity_(0) {}

    ...
    inline int num() const { return num_; }
    inline int channels() const { return channels_; }
    inline int height() const { return height_; }
    inline int width() const { return width_; }
    inline int count() const { return count_; }
    ...
    const Dtype* cpu_data() const;
    void set_cpu_data(Dtype* data);
    const Dtype* gpu_data() const;
    const Dtype* cpu_diff() const;
    const Dtype* gpu_diff() const;
    Dtype* mutable_cpu_data();
    Dtype* mutable_gpu_data();
    Dtype* mutable_cpu_diff();
    Dtype* mutable_gpu_diff();
    ...
}; // class Blob
```

Example - Layers

Protocol Buffers

```
message LayerParameter {
  repeated string bottom = 2; // the name of the bottom blobs
  repeated string top = 3; // the name of the top blobs
  optional string name = 4; // the layer name
  repeated NetStateRule include = 32;
  repeated NetStateRule exclude = 33;
  enum LayerType {
    ...
    CONVOLUTION = 4;
    ...
  }
  optional LayerType type = 5; // the layer type from the enum above

  repeated BlobProto blobs = 6;
  repeated float blobs_lr = 7;
  repeated float weight_decay = 8;
  repeated float loss_weight = 35;
  ...
  optional ContrastiveLossParameter contrastive_loss_param = 40;
  ...
  optional TransformationParameter transform_param = 36;
}
```

Example - Layers


- C++ class implementations
 - Setup: initialize the layer and connections
 - Forward: given input from bottoms compute outputs and pass to tops
 - Backward: given gradients wrt to top outputs, compute gradients wrt to input and sent to bottom. A layer with parameters computes gradients wrt to its parameters and store then internally
- A layer can have both CPU and GPU implementations for forward and backward propagations

Caffe Net

```
name: "CaffeNet"
layers {
  name: "data"
  type: IMAGE_DATA
  ...
}
layers {
  name: "conv1"
  type: CONVOLUTION
  ...
}
...

layers {
  name: "loss"
  type: SOFTMAX_LOSS
  ...
}
```

```
message NetParameter {
  optional string name = 1;
  repeated LayerParameter layers =
    2;
  repeated string input = 3;
  repeated int32 input_dim = 4;
  optional bool force_backward = 5
    [default = false];
  optional NetState state = 6;
}
```



Text
Representations



Net Definitions

Parameter
Representations



Data Structure

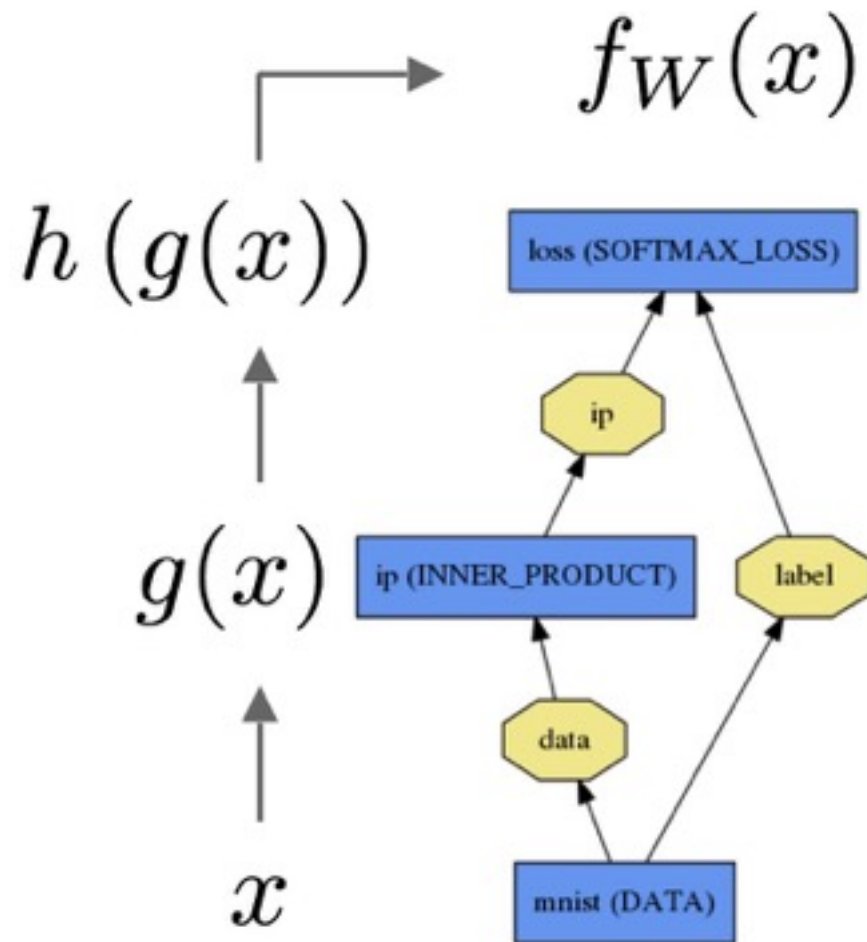
Class
Representations



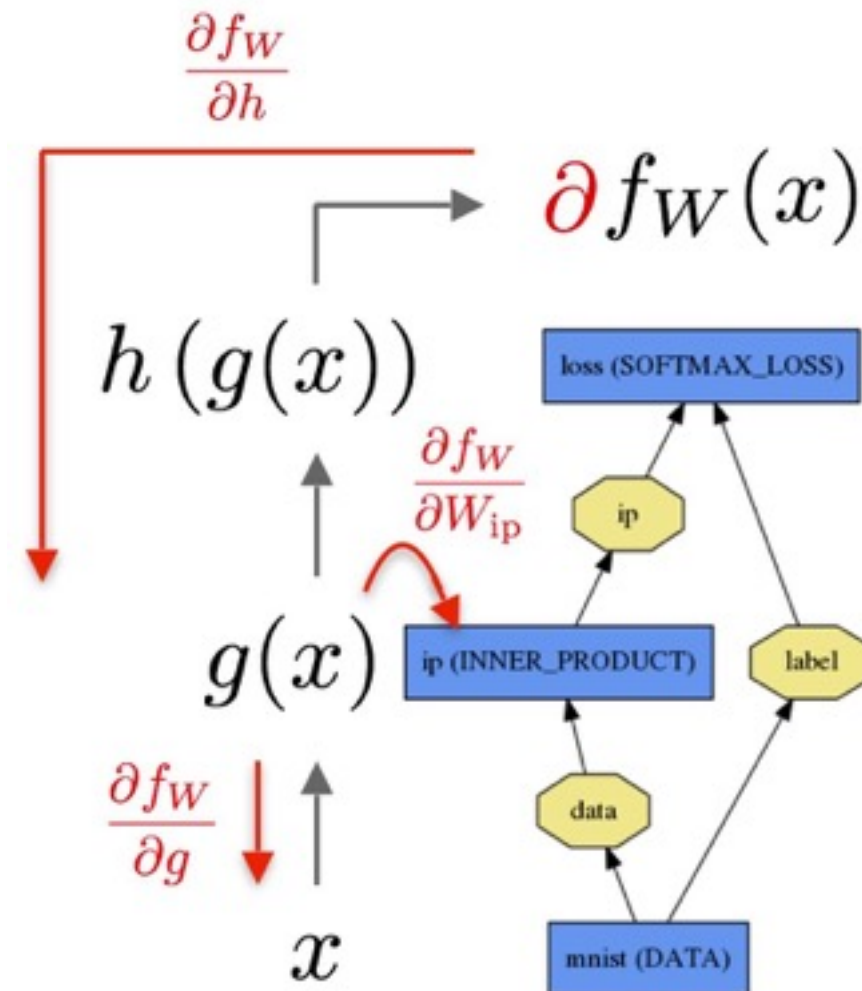
Implementation

Forward and Backward Propagation

Forward and Backward



```
Net::Forward();
Layer::Forward();
```



```
Net::Backward();
Layer::Backward();
```

Forward and Backward

- Each layer has CPU and GPU implementations for forward and backward:
 - Forward_cpu(), Forward_gpu()
 - Backward_cpu(), Backward_gpu()
- Solver calls forward first to get output and loss, then calls backward to calculate gradients of the model and update parameters

Loss and Supervision

Loss and Supervision

- Caffe is mainly used in supervised learning
- A loss function is the target function that we want to minimize
- Built-in loss functions:
 - Euclidean loss
 - Hinge loss
 - Multinomial logistic regression loss
 - Sigmoid cross entropy loss
 - Softmax loss
 - ...
- Multiple loss functions for multitasking

Loss and Supervision

Single loss function

```
layers {  
  name: "loss"  
  type: SOFTMAX_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

Multiple loss functions with weights

```
layers {  
  name: "loss_softmax"  
  type: SOFTMAX_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss_softmax"  
  loss_weight: 1  
}
```

```
layers {  
  name: "loss_eclidean"  
  type: ECLIDEAN_LOSS  
  bottom: "pred"  
  bottom: "label"  
  top: "loss_eclidean"  
  loss_weight: 0.5  
}
```

Caffe sums up all losses with non-zero loss_weight

Solver and optimizations

Solver and optimizations

- Solver
 - Creates training and validation networks
 - Iteratively calls forward/backward and updates parameters
 - Evaluates validation networks periodically
 - Snapshots model and solver states periodically

Solver and optimizations

Solver example

```
net: "examples/hdf5_classification/train_val.prototxt"  
test_iter: 1000  
test_interval: 1000  
base_lr: 0.01  
lr_policy: "step"  
gamma: 0.1  
stepsize: 5000  
display: 1000  
max_iter: 10000  
momentum: 0.9  
weight_decay: 0.0005  
snapshot: 10000  
snapshot_prefix: "examples/hdf5_classification/data/train"  
solver_mode: GPU
```

Data and Database

Data and Database

- Data are represented in Caffe as Blobs
- Data layers provide data to the network
 - Data can also be represented in Protocol Buffers
 - Structured data are serialized in binary format and stored in databases
- Databases
 - LevelDB
 - LMDB
 - HDF5
 - ...

Data Layers

- Tops and bottoms: multiple tops (data, label, ...), no bottoms
- Transformations: scaling, extracting mean, cropping, mirroring
- Prefetching: loading data in another thread
- Built-in data layers: DataLayer, ImageDataLayer, WindowDataLayer, HDF5DataLayer,...
- Multiple data layers or custom data layers for specific applications

Example - DataLayer

```
layers {  
  name: "data"  
  type: DATA  
  top: "data"  
  top: "label"  
  data_param {  
    source: "examples/imagenet/ilsvrc12_train_leveldb"  
    batch_size: 256  
  }  
  transform_param {  
    crop_size: 227  
    mean_file: "data/ilsvrc12/imagenet_mean.binaryproto"  
    mirror: true  
  }  
  include: { phase: TRAIN }  
}
```

Example - ImageDataLayer

```
layers {  
  name: "data"  
  type: IMAGE_DATA  
  top: "data"  
  top: "label"  
  image_data_param {  
    source: "examples/_temp/file_list.txt"  
    batch_size: 50  
    new_height: 256  
    new_width: 256  
  }  
  transform_param {  
    crop_size: 227  
    mean_file: "data/ilsvrc12/imagenet_mean.binaryproto"  
    mirror: false  
  }  
}
```

Interfaces and Tools

Interfaces

- Command-line interfaces
- Python interfaces
- Matlab interfaces

Command-line Interfaces

Training

```
# train LeNet
caffe train -solver examples/mnist/lenet_solver.prototxt
# train on GPU 2
caffe train -solver examples/mnist/lenet_solver.prototxt -gpu 2
# resume training from the half-way point snapshot
caffe train -solver examples/mnist/lenet_solver.prototxt \
-snapshot examples/mnist/lenet_iter_5000.solverstate
```

Fine-tuning

```
# fine-tune CaffeNet model weights for style recognition
caffe train -solver examples/solver.prototxt \
-weights models/bvlc_reference_caffenet.caffemodel
```

Testing

```
# testing LeNet in lenet_train_test.prototxt
caffe test -model examples/mnist/lenet_train_test.prototxt \
-weights examples/mnist/lenet_iter_10000.caffemodel -gpu 0 \
-iterations 100
```

Tools

- `convert_imageset.bin`: convert image datasets
- `compute_image_mean.bin`: compute mean values
- `extract_features.bin`: extract features using trained models
- `parse_log.sh`: parse training log files
- ...

Developing New Layers in Caffe

Kai KANG
kkang@ee.cuhk.edu.hk

Outline

- Look before You Leap
- Dive into Development
- Test Your Layers

Look Before You Leap

Look Before You Jump

- Do you really need a new layer?
- Have others implemented it yet?
- Is it worth the effort?

Do You Need a New Layer?

- Caffe may already have the layer
 - Read the documentations (<http://caffe.berkeleyvision.org/doxygen/annotated.html>)
- You may combine existing layers to have same effect
- Discover special usages of existing layers to simply the problem

Example - Sum Feature Maps

- Use 1x1 convolutions
- Weights initialized to 1
- Bias initialized to 0
- Learning rates and weight decays set to 0

```
layer {  
  name: "fusion"  
  type: "Convolution"  
  bottom: "feature_maps"  
  top: "fusion"  
  param { lr_mult: 0 decay_mult: 0 }  
  param { lr_mult: 0 decay_mult: 0 }  
  convolution_param {  
    num_output: 1  
    pad: 0  
    kernel_size: 1  
    stride: 1  
    weight_filler { type: "constant" value: 1 }  
    bias_filler { type: "constant" value: 0 }  
  }  
}
```


Example - Binary Error

- Threshold layer to get decisions
- Elementwise layer to calculate difference
- Absolute layer to get error
- Write new layer to sum all elements

```
layer {  
  name: "fusion"  
  type: "Convolution"  
  bottom: "feature_maps"  
  top: "fusion"  
  param { lr_mult: 0 decay_mult: 0 }  
  param { lr_mult: 0 decay_mult: 0 }  
  convolution_param {  
    num_output: 1  
    pad: 0  
    kernel_size: 1  
    stride: 1  
    weight_filler { type: "constant" value: 1 }  
    bias_filler { type: "constant" value: 0 }  
  }  
}
```

Have Others Done It?

- Discuss with other users
 - Colleagues or labmates
 - Caffe-user group
- Search the web
 - Search Issues or Pull-request on GitHub
 - Search related papers and ask for code
 - Search other users repositories

Example - Sum All Elements

BVLC / caffe

Watch 507

Unstar 2,747

Fork 1,554

Code 34

Issues 47

States

Closed 35

Open 12

Search all of GitHub

sum

Search

We've found 47 issues

Sort: Best match



ReductionLayer

#2089

Performs a "reduction" operation (currently **SUM**, MEAN, ASUM (**sum** of absolute values), and **SUM_OF_SQUARES**) to turn a number of "tail" axes into a single scalar value. The MEAN operation, in ...

Opened by jeffdonahue 7 days ago



MVNLayr across_channels is broken

#2017

Opened by cdoersch 14 days ago 3 comments



implementation of local fully connected layer

#1875

@ducha-aiki Thanks for the hint. I checked #1271, but it looks like not what I want. For example, the input is 6x6x256, the local connected layer in #1271 has a different weighted **sum** on the 256 ...

Opened by metalbubble 29 days ago 2 comments




Add parameter layer for learning any bottom

#2079

Opened by longjon 8 days ago 1 comment

Example - Sum All Elements

ReductionLayer #2089



jeffdonahue wants to merge 1 commit into `BVLC:master` from `jeffdonahue:reduction-layer`



Conversation 0



Commits 1



Files changed 5



Change Details

+508 -1



jeffdonahue commented 7 days ago

Owner

Performs a "reduction" operation (currently SUM, MEAN, ASUM (sum of absolute values), and SUM_OF_SQUARES) to turn a number of "tail" axes into a single scalar value. The `MEAN` operation, in combination with a `loss_weight`, is useful for creating custom losses that don't have an obnoxious amount of output. For example, this `EuclideanLoss`:

```
layer {
  type: "EuclideanLoss"
  bottom: "preds"
  bottom: "targets"
  top: "mean_squared_error"
  loss_weight: 1
}
```

...is equivalent to this `Reduction`:

```
layer {
  type: "Eltwise"
  bottom: "preds"
  bottom: "targets"
```

Labels

ready for review

Milestone

No milestone

Assignee

No one assigned

Notifications

Unsubscribe

You are receiving notifications because you're subscribed to this thread.

1 participant

Subscribe to get updates

Is It Worth the Effort?

- Writing new layers takes time, especially for the first time
- Maybe other frameworks are better options
 - Unsupervised training
 - Generative models

Dive into Development

Dive into Development

- Choose a good starting point
- Find a reference layer
- Follow the development procedure

Starting Point

- Clone the latest Caffe repository from GitHub
 - Avoid known issues
 - Newer versions make development easier
 - Easier to maintain for the future
- Compile and run the test first
- Checkout a new branch and start writing

```
# Clone the latest repository from
GitHub
# https://github.com/BVLC/caffe
git clone git@github.com:BVLC/caffe.git
```

```
# Compile and run the test
# http://caffe.berkeleyvision.org/
installation.html
make all
make test
make runtest
```

```
# Checkout a new branch and start
writing
git checkout -b mydev
```


Reference Layer

- Layer types
 - Common layers: ArgMax, ConCat, Eltwise, Flatten, InnerProduct, Silence, Softmax...
 - Data layers: Data, ImageData, WindowData...
 - Loss layers: Accuracy, EuclideanLoss, HingeLoss, SigmoidCrossEntropy...
 - Vision layers: Convolution, Pooling, Deconvolution
 - Neuron layers: ReLU, Sigmoid, Tanh...

Development Procedure

- Development wiki (<https://github.com/BVLC/caffe/wiki/Development>) for latest guides
- Add class declarations
- Implement you layers in C++ and CUDA
- Add Protocol Buffers parameters
- Add layer instantiation and registration
- Compile and debug

1. Add Class Declaration

- Add a class declaration in one of `common_layer.hpp`, `data_layers.hpp`, `loss_layers.hpp`, `neuron_layers.hpp` or `vision_layers.hpp`
- Include an inline implementation of `type()` and `*Blobs()` methods
- Omit `*_gpu` declarations if you'll only implement CPU code

Example - MapDataLayer

data_layer.hpp


```
template<typename Dtype>
class MapDataLayer : public BasePrefetchingDataLayer<Dtype> {
public:
    explicit MapDataLayer(const LayerParameter& param)
        : BasePrefetchingDataLayer<Dtype>(param),
          label_transformer_(label_trans_param(param.transform_param()), this->phase_) {}
    virtual ~MapDataLayer();
    virtual void DataLayerSetUp(const vector<Blob<Dtype>*>& bottom,
                                const vector<Blob<Dtype>*>& top);

    virtual inline const char* type() const { return "MapData"; }
    virtual inline int ExactNumBottomBlobs() const { return 0; }
    virtual inline int ExactNumTopBlobs() const { return 2; }

protected:
    virtual void InternalThreadEntry();
    DataTransformer<Dtype> label_transformer_;
    Blob<Dtype> transformed_label_;

    shared_ptr<db::DB> db_;
    shared_ptr<db::Cursor> iter_;

private:
    static TransformationParameter label_trans_param(
        const TransformationParameter& trans_param);
};
```



type()

*Blobs()

2. Implementation

- Implement your layer in `layers/your_layer.cpp`
 - (optional) `LayerSetUp` method for one-time initialization
 - `Reshape` method for computing sizes of top blobs, allocating buffers
 - `Forward_cpu` for forward propagation
 - `Backward_cpu` for gradient computations if needed
- Implement GPU versions `Forward_gpu` and `Backward_gpu` in `layers/your_layer.cu`

Example - MapDataLayer

map_data_layer.cpp

```
template <typename Dtype>
void MapDataLayer<Dtype>::DataLayerSetUp(const vector<Blob<Dtype>*>& bottom,
    const vector<Blob<Dtype>*>& top) {
    // Initialize DB
    // Read a data point and use it to initialize the top blob.
    // reshape data map
}

template<typename Dtype>
void MapDataLayer<Dtype>::InternalThreadEntry() {
    // Variable declarations
    // ...

    for (int item_id = 0; item_id < batch_size; ++item_id) {
        maps.ParseFromString(iter_->value());
        // Apply data and label transformations (mirror, scale, crop...)

        // go to the next iter
        iter_->Next();
        if (!iter_->valid()) {
            iter_->SeekToFirst();
        }
    }
}

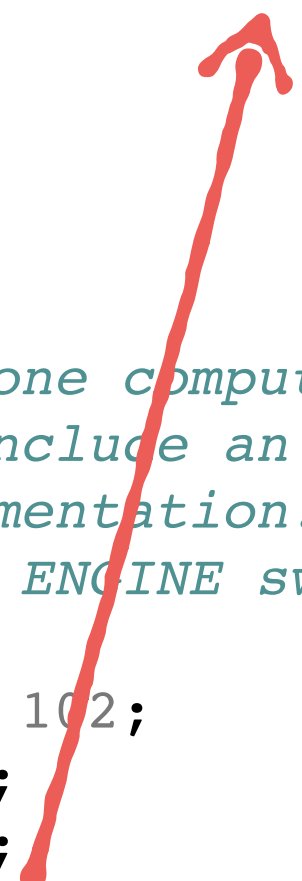
}
```

← Forward() for PrefetchDataLayer

3. Proto Declarations

- If needed, declare parameters in proto/caffe.proto

```
// NOTE
// Update the next available ID when you add a new LayerParameter field.
//
// LayerParameter next available layer-specific ID: 135 (last added:
reduction_param)
message LayerParameter {
    ...
    // Layer type-specific parameters.
    //
    // Note: certain layers may have more than one computational engine
    // for their implementation. These layers include an Engine type and
    // engine parameter for selecting the implementation.
    // The default for the engine is set by the ENGINE switch at compile-
    time.
    optional AccuracyParameter accuracy_param = 102;
    optional ArgMaxParameter argmax_param = 103;
    optional ConcatParameter concat_param = 104;
    optional MyLayerParameter mylayer_param = 135;
    ...
}
```



3. Proto Declarations

- If needed, declare parameters in proto/caffe.proto

```
// Message that stores parameters used by EltwiseLayer
message EltwiseParameter {
    enum EltwiseOp {
        PROD = 0;
        SUM = 1;
        MAX = 2;
    }
    optional EltwiseOp operation = 1 [default = SUM]; // element-wise
operation
    repeated float coeff = 2; // blob-wise coefficient for SUM operation

    // Whether to use an asymptotically slower (for >2 inputs) but stabler
method
    // of computing the gradient for the PROD operation. (No effect for SUM
op.)
    optional bool stable_prod_grad = 3 [default = true];
}
```


3. Layer Instantiation and Registration

- Add layer instantiation and registration in `my_layer.cpp`

```
INstantiate_Class(MyLayer);  
REGISTER_LAYER_CLASS(My);
```

- And in `my_layer.cu`

```
INstantiate_Layer_GPU_Funcs(MyLayer);
```

Layer Testing and Maintenance

Layer Testing and Maintenance

- Write tests in `test/test_your_layer.cpp`.
- Use `test/test_gradient_check_util.hpp` to check Forward and Backward implementations are in numerical agreement
- Write tests in different scenarios
- Compile, run tests and fix bugs
- Frequently pull the latest changes and update your own layers if there are conflicts

ELEG 4040

Tutorial 9

Caffe Reference Models

Kai KANG
kkang@ee.cuhk.edu.hk


Reference Models

- AlexNet
- GoogLeNet
- Network in Network
- FCN-Xs
- ...

AlexNet

- ImageNet 2012 Image Classification Challenge winner
- Originally trained on 2 GTX 580 GPUs because of insufficient memory
- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

AlexNet




```

layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96 kernel_size: 11 stride: 4
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0 }
  }
}

```

put all 2x48 channels on a single GPU

AlexNet




Local Response Normalization Layer

```

layer {
  name: "relu1" type: "ReLU" bottom: "conv1" top: "conv1"
}
layer {
  name: "norm1" type: "LRN" bottom: "conv1" top: "norm1"
  lrn_param { local_size: 5 alpha: 0.0001 beta: 0.75 }
}
layer {
  name: "pool1" type: "Pooling" bottom: "norm1" top: "pool1"
  pooling_param { pool: MAX kernel_size: 3 stride: 2 }
}

```


AlexNet




```

layer {
  name: "conv2"
  type: "Convolution"
  bottom: "pool1"
  top: "conv2"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 256 pad: 2 kernel_size: 5 group: 2
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0.1 }
  }
}

```

Divide input and output into two groups,
equivalent to put on two GPUs


AlexNet



```
layer {
  name: "conv3"
  type: "Convolution"
  bottom: "pool2"
  top: "conv3"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 384 pad: 1 kernel_size: 3
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0 }
  }
}
```

No groups, use all input channels

AlexNet




```

layer {
  name: "conv4"
  type: "Convolution"
  bottom: "conv3"
  top: "conv4"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 384 pad: 1 kernel_size: 3 group: 2
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0.1 }
  }
}
    
```

Divide input and output into two groups,
equivalent to put on two GPUs


AlexNet



```
layer {
  name: "conv5"
  type: "Convolution"
  bottom: "conv4"
  top: "conv5"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 256 pad: 1 kernel_size: 3 group: 2
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0.1 }
  }
}
```

Divide input and output into two groups,
equivalent to put on two GPUs

AlexNet




```

layer {
  name: "fc6"
  type: "InnerProduct"
  bottom: "pool5"
  top: "fc6"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  inner_product_param {
    num_output: 4096
    weight_filler { type: "gaussian" std: 0.005 }
    bias_filler { type: "constant" value: 0.1 }
  }
}

```

AlexNet




```

layer {
  name: "relu6"
  type: "ReLU"
  bottom: "fc6"
  top: "fc6"
}
layer {
  name: "drop6"
  type: "Dropout"
  bottom: "fc6"
  top: "fc6"
  dropout_param { dropout_ratio: 0.5 }
}

```

AlexNet




```

layer {
  name: "fc8"
  type: "InnerProduct"
  bottom: "fc7"
  top: "fc8"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  inner_product_param {
    num_output: 1000
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0 }
  }
}

```

1000 neurons for
1000 classes

AlexNet



```
layer {  
  name: "loss"  
  type: "SoftmaxWithLoss"  
  bottom: "fc8"  
  bottom: "label"  
  top: "loss"  
}
```


AlexNet

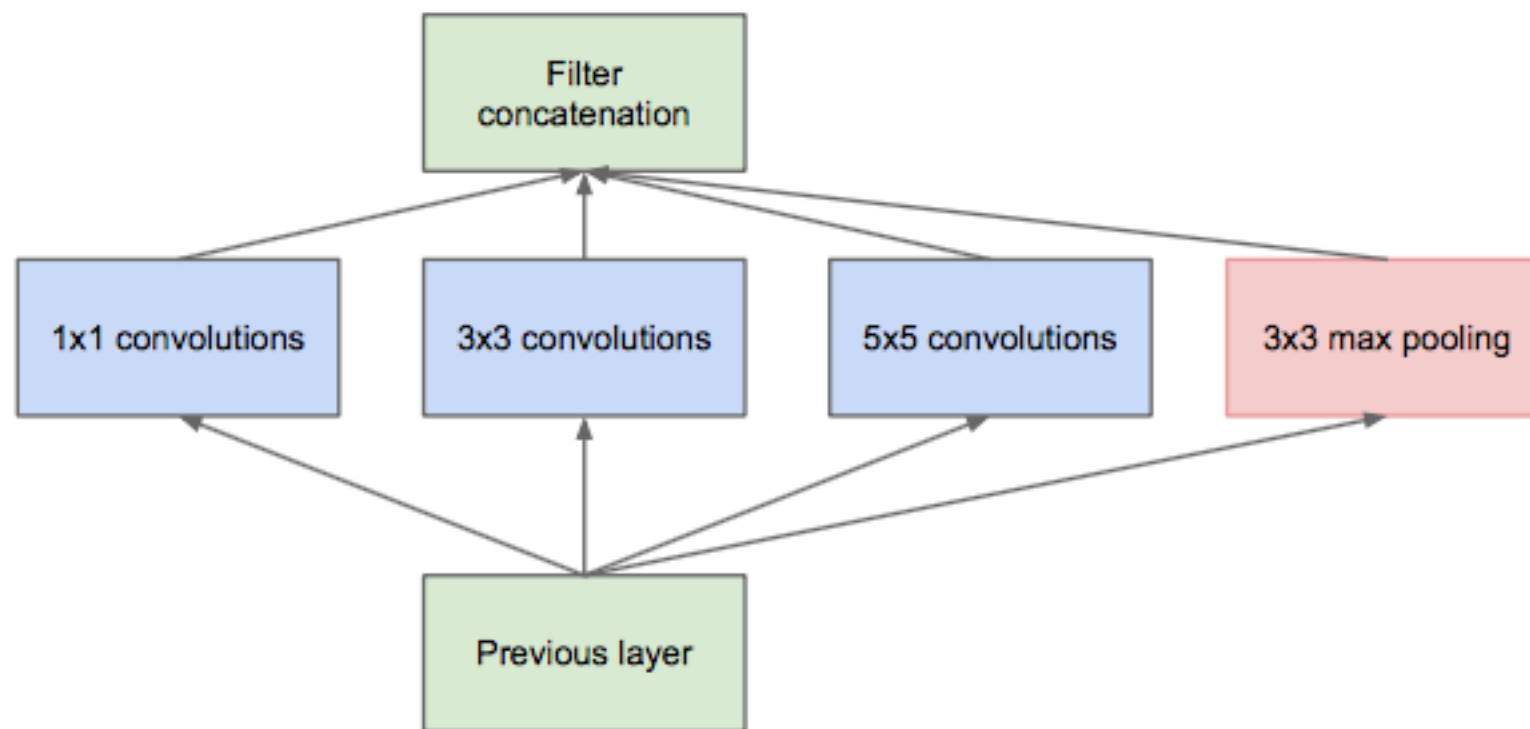
- Use **group** to convolve part of input channels to simulate two-GPU implementations
- Use local response normalization layer to normalize input responses
- Differences to original implementation
 - no relighting data augmentation
 - initializing non-zero biases to 0.1 instead of 1

GoogLeNet

- ImageNet 2014 Image Classification and Object Detection winner
- Szegedy, Christian, et al. "Going deeper with convolutions." arXiv preprint arXiv:1409.4842 (2014).

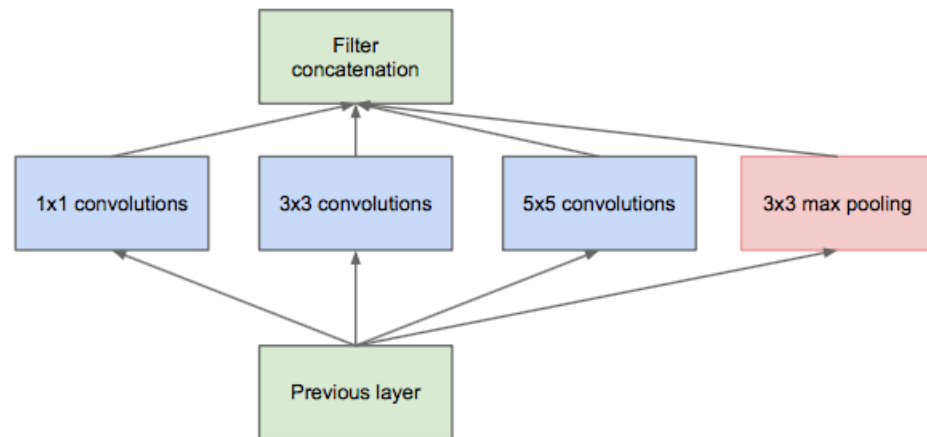
GoogLeNet

Inception Module - Naive Version



GoogLeNet

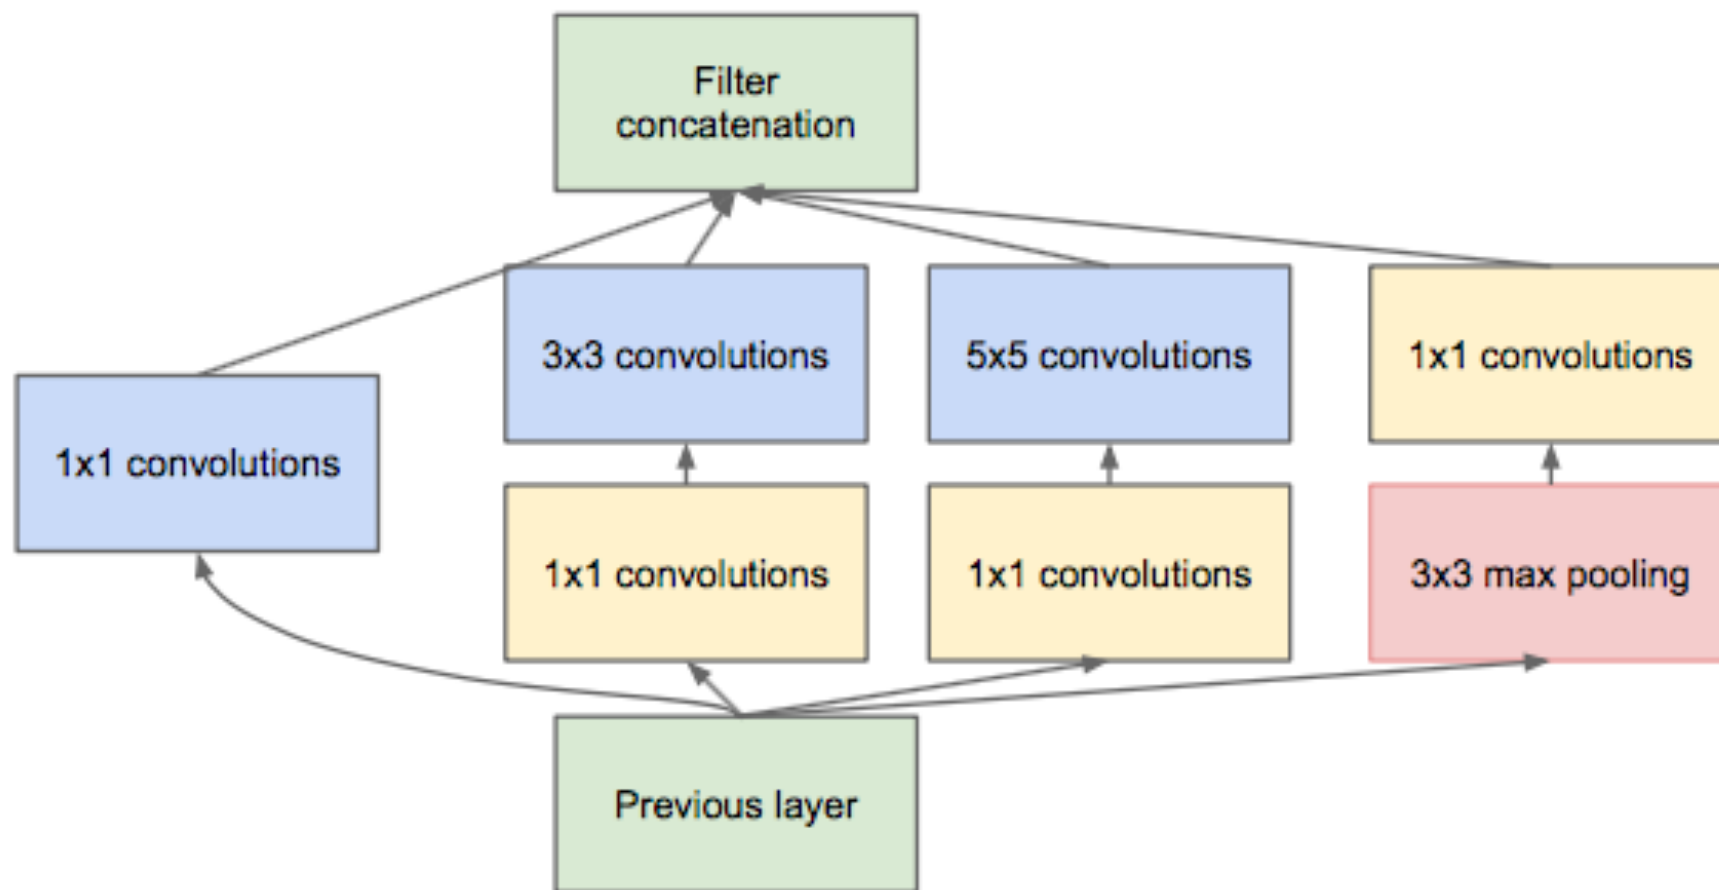
Inception Module - Naive Version



```
layer { name: "inception/1x1" type: "Convolution"
  bottom: "previous" top: "inception/1x1" ... }
layer { name: "inception/relu_1x1" type: "ReLU"
  bottom: "inception/1x1" top: "inception/1x1" }
layer { name: "inception/3x3" type: "Convolution"
  bottom: "previous" top: "inception/3x3" ... }
layer { name: "inception/relu_3x3" type: "ReLU"
  bottom: "inception/3x3" top: "inception/3x3" }
layer { name: "inception/5x5" type: "Convolution"
  bottom: "previous" top: "inception/5x5" ... }
layer { name: "inception/relu_5x5" type: "ReLU"
  bottom: "inception/5x5" top: "inception/5x5" ... }
layer { name: "inception/pool" type: "Pooling"
  bottom: "previous" top: "inception/pool"
  pooling_param { pool: MAX kernel_size: 3 stride: 1 pad: 1 } }
layer { name: "inception/relu_pool_proj" type: "ReLU"
  bottom: "inception/pool_proj" top: "inception/pool_proj" ... }
layer { name: "inception/output"
  type: "Concat"
  bottom: "inception/1x1"
  bottom: "inception/3x3"
  bottom: "inception/5x5"
  bottom: "inception/pool_proj"
  top: "inception/output"
}
```

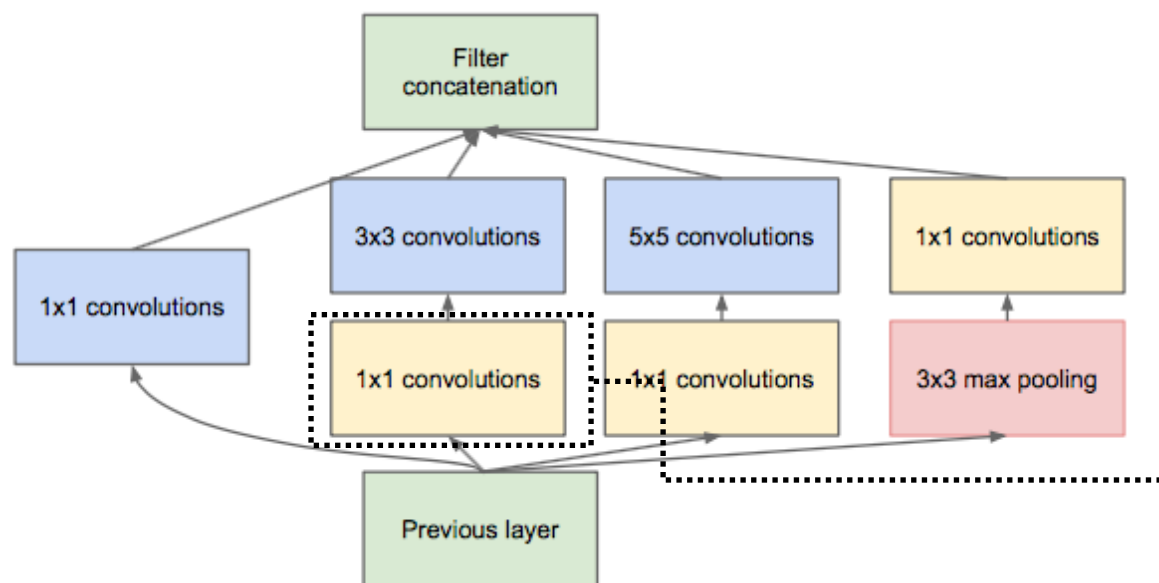
GoogLeNet

Inception Module - Dimension Reduction



GoogLeNet

Inception Module - Dimension Reduction



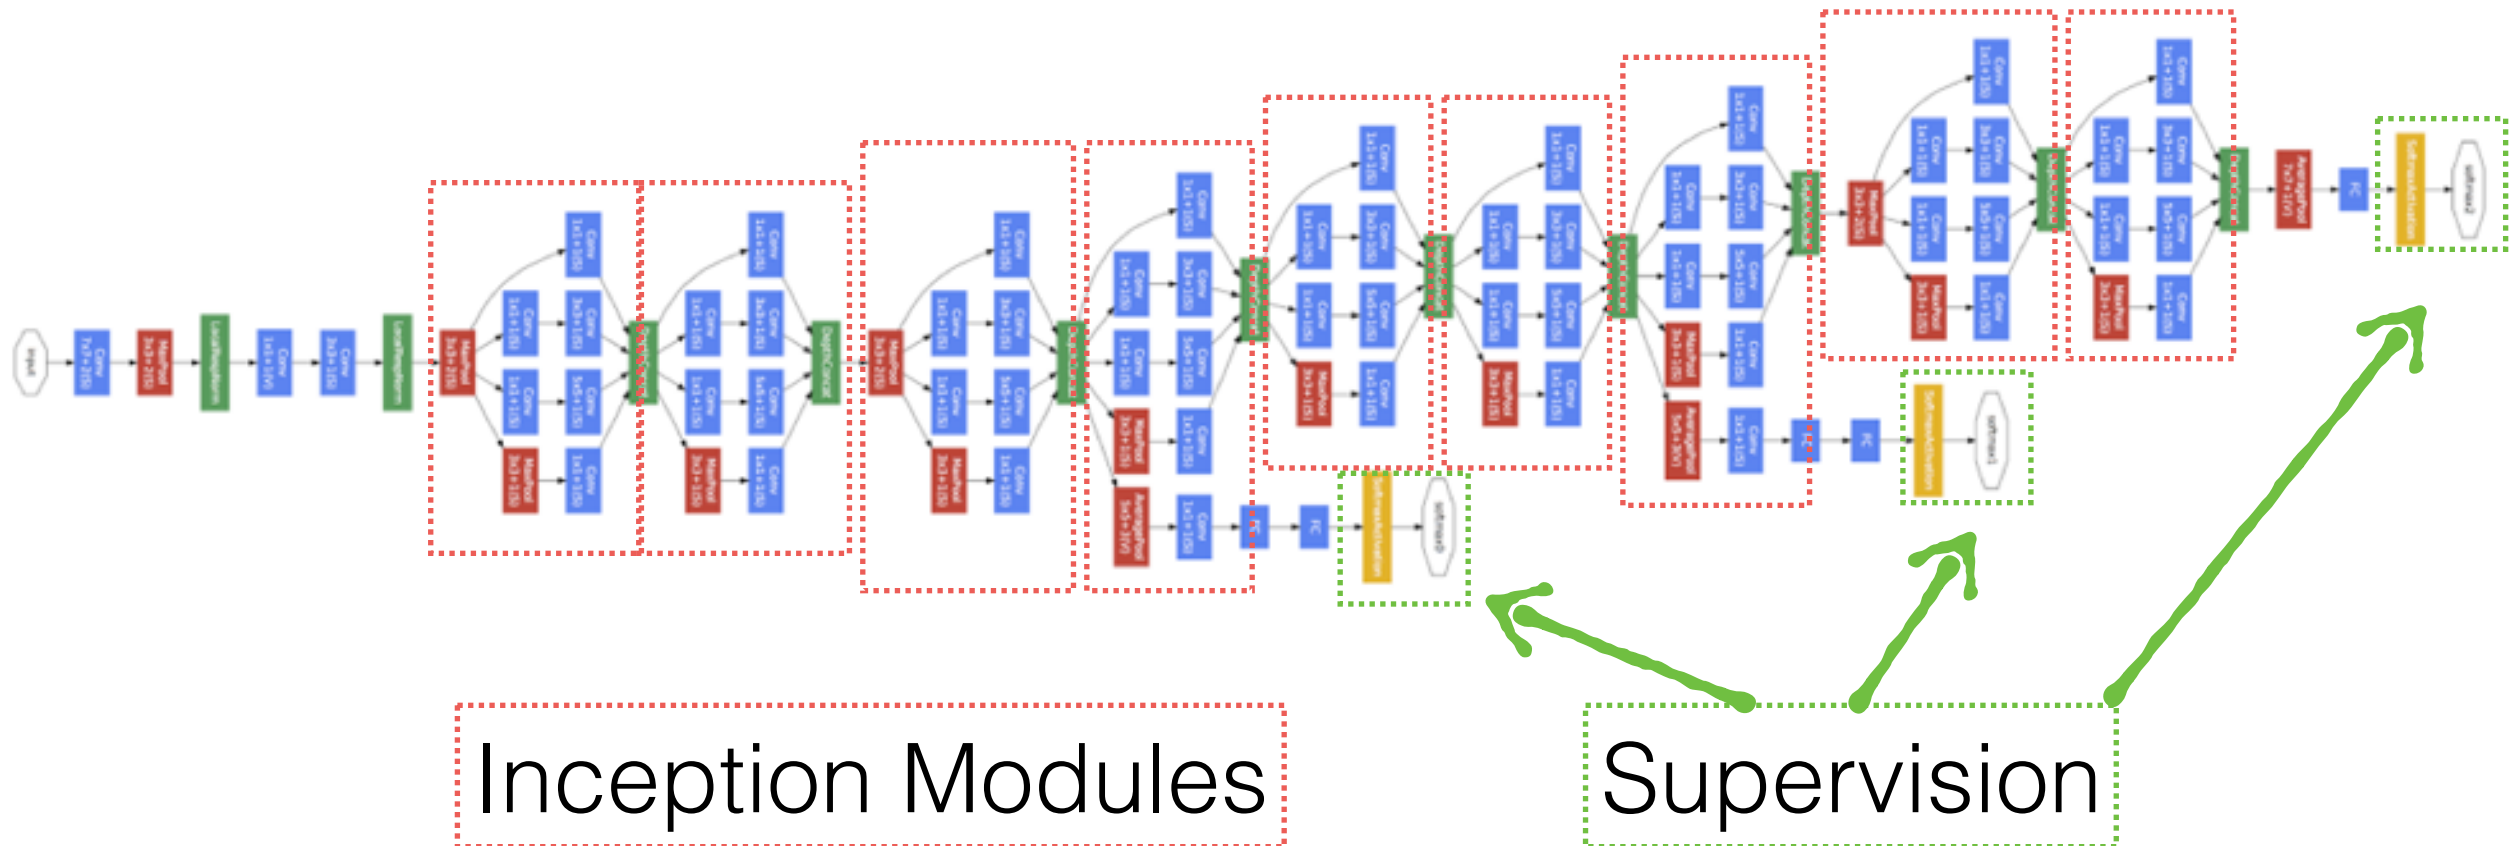
```
layer {  
  name: "inception/3x3_reduce"  
  type: "Convolution"  
  bottom: "previous"  
  top: "inception/3x3_reduce"  
  param { lr_mult: 1 decay_mult: 1 }  
  param { lr_mult: 2 decay_mult: 0 }  
  convolution_param {  
    num_output: 96 kernel_size: 1  
    weight_filler { type: "xavier" std: 0.09 }  
    bias_filler { type: "constant" value: 0.2 } }  
}  
layer {  
  name: "inception/relu_3x3_reduce"  
  type: "ReLU"  
  bottom: "inception/3x3_reduce"  
  top: "inception/3x3_reduce"  
}
```

reduce channels with 1x1 kernels

Net Structure



GoogLeNet



Inception Modules

Supervision

```
layer {
  name: "loss1/classifier"
  type: "InnerProduct"
  bottom: "loss1/fc"
  top: "loss1/classifier"
  param { lr_mult: 1 decay_mult: 1 }
  param { lr_mult: 2 decay_mult: 0 }
  inner_product_param {
    num_output: 1000
    weight_filler { type: "xavier" std: 0.0009765625 }
    bias_filler { type: "constant" value: 0 }
  }
}
```

```
layer {
  name: "loss1/loss"
  type: "SoftmaxWithLoss"
  bottom: "loss1/classifier"
  bottom: "label"
  top: "loss1/loss1"
  loss_weight: 0.3
}
```



GoogLeNet

- Differences:
 - not training with the relighting data-augmentation;
 - not training with the scale or aspect-ratio data-augmentation;
 - uses "xavier" to initialize the weights instead of "gaussian";

Network in Network

- ImageNet 2014 Object Detection (with provided data only) winner
- Use global pooling rather than fully connected layers for classification
- Lin, Min, Qiang Chen, and Shuicheng Yan. "Network in network." arXiv preprint arXiv:1312.4400 (2013).


Network in Network



```
layers {  
  bottom: "conv1"  
  top: "cccp1"  
  name: "cccp1"  
  type: CONVOLUTION  
  convolution_param {  
    num_output: 96  
    kernel_size: 1  
    stride: 1  
    weight_filler { type: "gaussian" mean: 0 std: 0.05 }  
    bias_filler { type: "constant" value: 0 }  
  }  
}  
layers {  
  bottom: "cccp1"  
  top: "cccp1"  
  name: "relu1"  
  type: RELU  
}
```

```
layers {  
  bottom: "cccp1"  
  top: "cccp2"  
  name: "cccp2"  
  type: CONVOLUTION  
  convolution_param {  
    num_output: 96  
    kernel_size: 1  
    stride: 1  
    weight_filler { type: "gaussian" mean: 0 std: 0.05 }  
    bias_filler { type: "constant" value: 0 }  
  }  
}  
layers {  
  bottom: "cccp2"  
  top: "cccp2"  
  name: "relu2"  
  type: RELU  
}
```

Network in Network




```
layers {  
  bottom: "cccp8"  
  top: "pool4"  
  name: "pool4"  
  type: POOLING  
  pooling_param {  
    pool: AVE  
    global_pooling: true  
  }  
}
```

FCN-Xs

- State-of-the-art results on PASCAL VOC segmentation challenges
- Use trainable deconvolution layer to get dense prediction
- <https://github.com/longjon/caffe/tree/future>
- Long, Jonathan, Evan Shelhamer, and Trevor Darrell. "Fully convolutional networks for semantic segmentation." arXiv preprint arXiv:1411.4038 (2014).

FCN-Xs



```
layers { type: DECONVOLUTION name: 'score2' bottom: 'score' top: 'score2'
  blobs_lr: 1 blobs_lr: 2 weight_decay: 1 weight_decay: 0
  convolution_param { kernel_size: 4 stride: 2 num_output: 21 } }
```

```
layers { type: DECONVOLUTION name: 'score4' bottom: 'score-fused'
  top: 'score4'
  blobs_lr: 1 weight_decay: 1
  convolution_param { bias_term: false kernel_size: 4 stride: 2 num_output: 21 } }
```

```
layers { type: DECONVOLUTION name: 'upsample'
  bottom: 'score-final' top: 'bigscore'
  blobs_lr: 0
  convolution_param { bias_term: false num_output: 21 kernel_size: 16 stride: 8 } }
```

deconvolution
for upsampling